

# COMP 540 Assignment #1

Yunda Jia  
Yu Wu

January 24, 2020

## 0 Background refresher(30 points)

- Plot the categorical distribution.

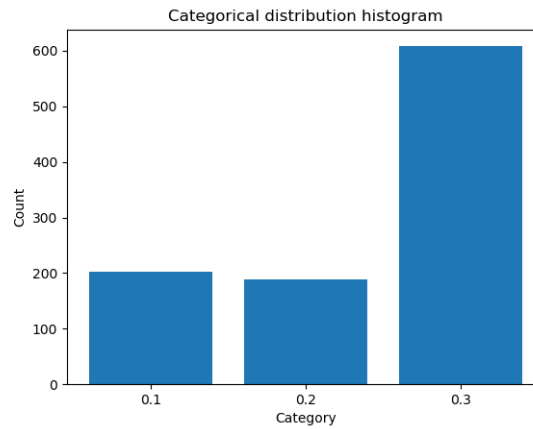


Figure 1: Categorical distribution

- Plot the Univariate normal distribution with mean of and standard deviation of 1.




Figure 2: Univariate Normal Distribution

- Produce a scatter plot of the samples for a 2-D Gaussian.

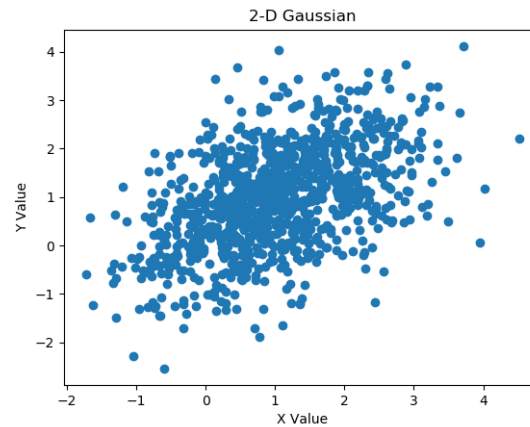


Figure 3: Univariate Normal Distribution

- Test mixture sampling code Code can be seen in `sampler.py`. Mixture Gaussian plot is shown below

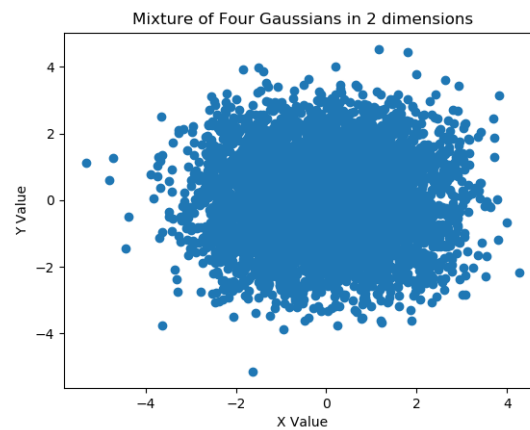


Figure 4: Univariate Normal Distribution

- Prove that the sum of two independent Poisson random variables is also a Poisson random variable. Suppose  $X \sim \mathcal{P}(\lambda)$  and  $Y \sim \mathcal{P}(\mu)$ . Now Prove that  $X + Y \sim \mathcal{P}(\lambda + \mu)$ .

$$\begin{aligned}
P(X + Y = k) &= \sum_{i=0}^k P(X + Y = k, X = i) \\
&= \sum_{i=0}^k P(Y = k - i, X = i) \\
&= \sum_{i=0}^k P(Y = k - i)P(X = i) \\
&= \sum_{i=0}^k e^{-\mu} \frac{\mu^{k-i}}{(k-i)!} e^{-\lambda} \frac{\lambda^i}{i!} \\
&= e^{-(\mu+\lambda)} \frac{1}{k!} \sum_{i=0}^k \frac{k!}{i!(k-i)!} \mu^{k-i} \lambda^i \\
&= e^{-(\mu+\lambda)} \frac{1}{k!} \sum_{i=0}^k \binom{k}{i} \mu^{k-i} \lambda^i \\
&= \frac{(\mu + \lambda)^k}{k!} \cdot e^{-(\mu+\lambda)}
\end{aligned}$$

So  $X + Y \sim \mathcal{P}(\lambda + \mu)$ .

- Find  $\alpha, \mu_1$  and  $\sigma_1$ .

We have  $X_0$  and  $X_1$  be continuous random variables. If

$$\begin{aligned}
p(X_0 = x_0) &= \alpha_0 e^{-\frac{(x_0 - \mu_0)^2}{2\sigma_0^2}} \\
P(X_1 = x_1 | X_0 = x_0) &= \alpha_1 e^{-\frac{(x_1 - x_0)^2}{2\sigma^2}}
\end{aligned}$$

$$\begin{aligned}
p(X_1 = x_1) &= \int P(X_1 = x_1 | X_0 = x_0) \cdot p(X_0 = x_0) dx_0 \\
&= \alpha_0 \alpha_1 \int e^{-\frac{\sigma^2(x_0 - \mu_0)^2 + \sigma_0^2(x_1 - x_0)^2}{2\sigma_0^2\sigma^2}} dx_0 \\
&= \alpha_0 \alpha_1 \int e^{-\frac{(\sigma^2 + \sigma_0^2)x_0^2 - 2(\sigma^2\mu_0 + \sigma_0^2x_1)x_0 + \sigma^2\mu_0^2 + \sigma_0^2x_1^2}{2\sigma_0^2\sigma^2}} dx_0 \\
&= \alpha_0 \alpha_1 \int e^{-\frac{\frac{1}{2\sigma_0^2\sigma^2}[(\sqrt{\sigma^2 + \sigma_0^2}x_0 - \frac{-\sigma^2\mu_0 + \sigma_0^2x_1}{\sqrt{\sigma^2 + \sigma_0^2}})^2 + \sigma^2\mu_0^2 + \sigma_0^2x_1^2 - \frac{\sigma^2\mu_0^2 + \sigma_0^2x_1^2}{\sigma^2 + \sigma_0^2}]}{2\sigma_0^2\sigma^2}} dx_0
\end{aligned}$$

Since

$$\int \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx = 1$$

$$\begin{aligned}
p(X_1 = x_1) &= \frac{\alpha\alpha_0\sqrt{2\pi}\sigma\sigma_0}{\sqrt{\sigma^2 + \sigma_0^2}} e^{-\frac{\frac{1}{2\sigma_0^2\sigma^2}(\sigma^2\mu_0^2 + \sigma_0^2x_1^2 - \frac{\sigma^2\mu_0^2 + \sigma_0^2x_1^2}{\sigma^2 + \sigma_0^2})}{2\sigma_0^2\sigma^2}} \\
&= \alpha e^{-\frac{(x_1 - \mu_0)^2}{2(\sigma^2 + \sigma_0^2)}}
\end{aligned}$$

Thus we can solve that:

$$\alpha = \frac{\alpha\alpha_0\sqrt{2\pi\sigma\sigma_0}}{\sqrt{\sigma^2 + \sigma_0^2}}$$

$$\mu_1 = \mu_0$$

$$\sigma_1 = \sqrt{\sigma^2 + \sigma_0^2}$$

- Show that if  $P(A|B, C) > P(A|B)$  then  $P(A|B, C^C) < P(A|B)$

$$P(A|B, C) = \frac{P(A \cap B \cap C)}{P(B \cap C)}$$

$$\frac{P(A \cap B \cap C)}{P(B \cap C)} > \frac{P(A \cap B)}{P(B)}$$

$$\frac{P(C|A \cap B)}{P(C|B)} > 1$$

$$P(C|A \cap B) > P(C|B)$$

$$1 - P(C|A \cap B) < 1 - P(C|B)$$

$$P(C^C|A \cap B) < P(C^C|B)$$

$$\frac{P(C^C|A \cap B)}{P(C^C|B)} < 1$$

$$\begin{aligned} P(A|B, C^C) &= \frac{P(A \cap B \cap C^C)}{P(B \cap C^C)} \\ &= \frac{P(C^C|A \cap B)P(A \cap B)}{P(C^C|B)P(B)} \\ &= \frac{P(C^C|A \cap B)}{P(C^C|B)P(B)} P(A|B) \\ &< P(A|B) \end{aligned}$$

Thus, if  $P(A|B, C) > P(A|B)$  then  $P(A|B, C^C) < P(A|B)$ .

- Consider the vectors  $\mathbf{u} = [1 \ 2]^T$  and  $\mathbf{v} = [2 \ 3]^T$ . Define the matrix  $\mathbf{M} = \mathbf{u}\mathbf{v}^T$ . Compute the eigenvalues and eigenvectors of  $\mathbf{M}$ .

$$\begin{aligned} \mathbf{M} &= \mathbf{u}\mathbf{v}^T \\ &= [1 \ 2]^T \cdot [2 \ 3] \\ &= \begin{bmatrix} 2 & 3 \\ 4 & 6 \end{bmatrix} \end{aligned}$$

$$|\lambda \mathbf{I} - \mathbf{M}| = \begin{vmatrix} \lambda - 2 & -3 \\ -4 & \lambda - 6 \end{vmatrix} = 0$$

Then we can solve it as

$$\begin{aligned}
(\lambda - 2)(\lambda - 6) - 12 &= 0 \\
\lambda^2 - 8\lambda &= 0 \\
\lambda(\lambda - 8) &= 0 \\
\lambda_1 = 0, \lambda_2 = 8
\end{aligned}$$

Let  $\lambda = 0$ :

$$\begin{aligned}
(\lambda \mathbf{I} - \mathbf{M}) \cdot \begin{bmatrix} x_1 & x_2 \end{bmatrix}^T &= 0 \\
\begin{bmatrix} -2 & -3 \\ -4 & -6 \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} &= \begin{bmatrix} 0 \\ 0 \end{bmatrix} \\
-2x_1 - 3x_2 &= 0
\end{aligned}$$

Let  $x_1 = 3$ , Then  $x_2 = -2$ . So eigenvector is  $\begin{bmatrix} 3 & -2 \end{bmatrix}^T$ .  
Let  $\lambda = 8$ :

$$\begin{aligned}
(\lambda \mathbf{I} - \mathbf{M}) \cdot \begin{bmatrix} x_1 & x_2 \end{bmatrix}^T &= 0 \\
\begin{bmatrix} 6 & -3 \\ -4 & 2 \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} &= \begin{bmatrix} 0 \\ 0 \end{bmatrix} \\
6x_1 - 3x_2 &= 0 \\
-4x_1 + 2x_2 &= 0
\end{aligned}$$

Let  $x_1 = 1$ , Then  $x_2 = 2$ . So eigenvector is  $\begin{bmatrix} 1 & 2 \end{bmatrix}^T$ .  
Thus, when eigenvalue  $\lambda = 0$ , eigenvector is  $\begin{bmatrix} 3 & -2 \end{bmatrix}^T$ , when eigenvalue  $\lambda = 8$ , eigenvector is  $\begin{bmatrix} 1 & 2 \end{bmatrix}^T$ .

- Show that if  $A$  is positive semi-definite, then all eigenvalues of  $A$  are non-negative.  
Definition of an eigenvalue and eigenvector is:

$$A\mathbf{v} = \lambda\mathbf{v}$$

If  $A$  is positive semi-definite,  $\mathbf{x}^T A \mathbf{x} \geq 0$  for all  $\mathbf{x}$ .

$$\mathbf{v}^T A \mathbf{v} = \mathbf{v}^T \mathbf{v} \lambda$$

Since  $\mathbf{v}^T \mathbf{v}$  is necessarily a positive number, in order for  $\mathbf{v}^T A \mathbf{v}$  to be greater than or equal to 0,  $\lambda$  must be  $\lambda \geq 0$ .

- Provide one example for each of the following cases.

As for  $(A + B)^2 \neq A^2 + 2AB + B^2$ . Suppose  $A = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$  and  $B = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$ . Since  $A \cdot B = \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix}$  and

$$B \cdot A = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}.$$

So  $(A + B)^2 = A^2 + AB + BA + B^2 \neq A^2 + 2AB + B^2$ .

As for  $AB = 0, A \neq 0, B \neq 0$ , Suppose  $A = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$  and  $B = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$ .

- Show that  $\mathbf{A}$  is orthogonal.  
Given  $\mathbf{u}^T \mathbf{u} = 1$  and  $\mathbf{A} = \mathbf{I} - 2\mathbf{u}\mathbf{u}^T$ .

$$\begin{aligned}
\mathbf{A}^T \mathbf{A} &= (\mathbf{I} - 2\mathbf{u}\mathbf{u}^T)^T (\mathbf{I} - 2\mathbf{u}\mathbf{u}^T) \\
&= (\mathbf{I} - 2\mathbf{u}\mathbf{u}^T)(\mathbf{I} - 2\mathbf{u}\mathbf{u}^T) \\
&= \mathbf{I} - 2\mathbf{u}\mathbf{u}^T - 2\mathbf{u}\mathbf{u}^T + 4\mathbf{u}\mathbf{u}^T \mathbf{u}\mathbf{u}^T \\
&= \mathbf{I} - 2\mathbf{u}\mathbf{u}^T - 2\mathbf{u}\mathbf{u}^T + 4\mathbf{u}\mathbf{u}^T \\
&= \mathbf{I}
\end{aligned}$$

So  $\mathbf{A}$  is orthogonal.

- Prove the following assertions.  
–As for  $f(x) = x^3$  for  $x \geq 0$ .

$$\begin{aligned}
f(x) &= x^3 \\
f''(x) &= 6x
\end{aligned}$$

Since  $x \geq 0$ ,  $f''(x) = 6x \geq 0$ . Thus  $f(x) = x^3$  is convex for  $x \geq 0$ .  
–As for  $f(x_1, x_2) = \max(x_1, x_2)$ . Let  $\mathbf{x} = (x_1, x_2)$  and  $\mathbf{y} = (y_1, y_2)$  and  $\lambda \in [0, 1]$ .

$$\begin{aligned}
f(\lambda \mathbf{x} + (1 - \lambda)\mathbf{y}) &= f(\lambda x_1 + (1 - \lambda)y_1, \lambda x_2 + (1 - \lambda)y_2) \\
&= \max(\lambda x_1 + (1 - \lambda)y_1, \lambda x_2 + (1 - \lambda)y_2) \\
&\leq \max(\lambda x_1, \lambda x_2) + \max((1 - \lambda)y_1, (1 - \lambda)y_2) \\
&= \lambda \max(x_1, x_2) + (1 - \lambda) \max(y_1, y_2) \\
&= \lambda f(\mathbf{x}) + (1 - \lambda)f(\mathbf{y})
\end{aligned}$$

So  $f(x_1, x_2) = \max(x_1, x_2)$  is convex on  $R^2$ .

–As for function  $f + g$ . If univariate functions  $f$  and  $g$  are convex on  $S$ , then  $f''(x) \geq 0$  and  $g''(x) \geq 0$ .

So  $(f + g)''(x) = f''(x) + g''(x) \geq 0$ . Thus if univariate functions  $f$  and  $g$  are convex on  $S$ , then  $f + g$  is convex on  $S$ .

–As for  $fg$ . If univariate functions  $f$  and  $g$  are convex and non-negative on  $S$ .

$$\begin{aligned}
(fg)''(x) &= (f'g + fg')'(x) \\
&= (f''g + f'g' + f'g' + fg'')(x) \\
&= (f''g + 2f'g' + fg'')(x)
\end{aligned}$$

Since  $f$  and  $g$  have their minimum within  $S$  at the same point. Before the minimum point, both  $f$  and  $g$  are decreasing. After the minimum point, both  $f$  and  $g$  are increasing. So  $f'g' \geq 0$ ,  $f'' \geq 0$ ,  $g'' \geq 0$ ,  $f \geq 0$  and  $g \geq 0$ . Thus  $(fg)''(x) = (f''g + 2f'g' + fg'')(x) \geq 0$ . Then  $fg$  is convex on  $S$ .

- Find the highest entropy of categorical distribution.  
The entropy of a categorical distribution on  $K$  values is defined as

$$H(p) = - \sum_{i=1}^K p_i \log(p_i)$$

The probability and constraints are defined below:

$$P(X = x_i) = p_i \quad \text{for } i = 1, 2, \dots, K$$

$$s.p. \begin{cases} \sum_{i=1}^K p_i = 1 \\ p_i \geq 0 \quad for \quad i = 1, 2, \dots, K \end{cases}$$

Constrain function is

$$\varphi(p_i) = \sum_{i=1}^K p_i - 1 = 0$$

Define Lagrange multipliers:

$$\begin{aligned} \mathcal{L} &= H(p) + \lambda \varphi(p_i) \\ &= -\sum_{i=1}^K p_i \log(p_i) + \lambda \left( \sum_{i=1}^K p_i - 1 \right) \end{aligned}$$

To find the highest entropy, we should find the point where derivative is 0.

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial p_i} &= -\log(p_i) - 1 + \lambda = 0 \\ p^* &= e^{\lambda-1} \end{aligned}$$

Thus when all  $p_i, i = 1, 2, \dots, K$  are equally equal to  $e^{\lambda-1}$ , the categorical distribution has the highest entropy.

## 1 Locally weighted linear regression(20 points)

- Find an appropriate diagonal matrix  $W$ .

$$J(\theta) = (X\theta - y)^T W (X\theta - y)$$

Let  $W$  be

$$W = \begin{bmatrix} \frac{1}{2}w^{(1)} & & & & 0 \\ & \frac{1}{2}w^{(2)} & & & \\ & 0 & & \ddots & \\ & & & & \frac{1}{2}w^{(i)} \end{bmatrix}$$

$X$  is the  $m \times d$  input matrix and  $y$  is a  $m \times 1$  vector.

$$\begin{aligned} X &= \begin{bmatrix} x_1^{(1)} & x_2^{(1)} & x_3^{(1)} & \cdots & x_d^{(1)} \\ x_1^{(2)} & x_2^{(2)} & x_3^{(2)} & \cdots & x_d^{(2)} \\ \vdots & \vdots & \vdots & \cdots & \vdots \\ x_1^{(m)} & x_2^{(m)} & x_3^{(m)} & \cdots & x_d^{(m)} \end{bmatrix} \\ y &= \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} \\ \theta &= \begin{bmatrix} \theta_{(1)} \\ \theta_{(2)} \\ \vdots \\ \theta_{(d)} \end{bmatrix} \end{aligned}$$

$$X\theta - y = \begin{bmatrix} (\theta_1 x_1^{(1)} + \theta_2 x_2^{(1)} + \dots + \theta_d x_d^{(1)}) - y^{(1)} \\ (\theta_1 x_1^{(2)} + \theta_2 x_2^{(2)} + \dots + \theta_d x_d^{(2)}) - y^{(2)} \\ \vdots \\ (\theta_1 x_1^{(m)} + \theta_2 x_2^{(m)} + \dots + \theta_d x_d^{(m)}) - y^{(m)} \end{bmatrix}$$

$$W(X\theta - y) = \begin{bmatrix} \frac{1}{2}w^{(1)} \times ((\theta_1 x_1^{(1)} + \theta_2 x_2^{(1)} + \dots + \theta_d x_d^{(1)}) - y^{(1)}) \\ \frac{1}{2}w^{(2)} \times ((\theta_1 x_1^{(2)} + \theta_2 x_2^{(2)} + \dots + \theta_d x_d^{(2)}) - y^{(2)}) \\ \vdots \\ \frac{1}{2}w^{(m)} \times ((\theta_1 x_1^{(m)} + \theta_2 x_2^{(m)} + \dots + \theta_d x_d^{(m)}) - y^{(m)}) \end{bmatrix}$$

$$\begin{aligned} (X\theta - y)^T W(X\theta - y) &= \frac{1}{2}w^{(1)} \times ((\theta_1 x_1^{(1)} + \theta_2 x_2^{(1)} + \dots + \theta_d x_d^{(1)}) - y^{(1)})^2 \\ &\quad + \frac{1}{2}w^{(2)} \times ((\theta_1 x_1^{(2)} + \theta_2 x_2^{(2)} + \dots + \theta_d x_d^{(2)}) - y^{(2)})^2 + \dots \\ &\quad + \frac{1}{2}w^{(m)} \times ((\theta_1 x_1^{(m)} + \theta_2 x_2^{(m)} + \dots + \theta_d x_d^{(m)}) - y^{(m)})^2 \end{aligned}$$

So  $J(\theta) = (X\theta - y)^T W(X\theta - y)$  can be written in the form  $J(\theta) = (X\theta - y)^T W(X\theta - y)$  when choosing  $W$  as above.

- Solve for  $\theta$

$$\begin{aligned} J(\theta) &= (X\theta - y)^T W(X\theta - y) \\ &= ((X\theta)^T - y^T) W(X\theta - y) \\ &= ((X\theta)^T W - y^T W)(X\theta - y) \\ &= \theta^T X^T W X\theta - 2\theta^T X^T W y + y^T W y \end{aligned}$$

$$\begin{aligned} \nabla_{\theta} J(\theta) &= 2X^T W X\theta - 2X^T W y = 0 \\ \theta &= (X^T W X)^{-1} X^T W y \\ &= (X^T X)^{-1} X^T y \end{aligned}$$

Thus  $\theta = (X^T X)^{-1} X^T y$ .

- Write down an algorithm for calculating  $\theta$  by batch gradient descent for locally weighted linear regression. We know that  $\theta = (X^T X)^{-1} X^T y$ .

Algorithm is

---

**Algorithm 1:** Calculating  $\theta$  by batch gradient descent

---

**Input:**  $x$

**Output:**  $\theta$

**for**  $i < 10000$  **do**

$\theta_j \leftarrow \theta_j - \alpha \sum_i w^{(i)} (\sum_k (\theta_k x_k^{(i)}) - y^{(i)}), j = 0$

$\theta_j \leftarrow \theta_j - \alpha \sum_i w^{(i)} (\sum_k (\theta_k x_k^{(i)}) - y^{(i)}) x_j^{(i)}, j \neq 0$

---

Locally weighted linear regression is a non-parametric method.



## 2 Properties of the linear regression estimator(10 points)

- Show that  $E[\theta] = \theta^*$

$$\begin{aligned}\theta &= (X^T X)^{-1} X^T y \\ &= (X^T X)^{-1} X^T (X\theta^* + \epsilon) \\ &= (X^T X)^{-1} X^T X\theta^* + (X^T X)^{-1} X^T \epsilon\end{aligned}$$

$$\begin{aligned}E(\theta) &= E((X^T X)^{-1} X^T X\theta^* + (X^T X)^{-1} X^T \epsilon) \\ &= E((X^T X)^{-1} X^T X\theta^*) + E((X^T X)^{-1} X^T \epsilon) \\ &= \theta^* + (X^T X)^{-1} X^T E(\epsilon) \\ &= \theta^*\end{aligned}$$

- Show the variance of the least squares estimator is  $Var(\theta) = (X^T X)^{-1} \sigma^2$ . Suppose  $\mathbf{b}$  is an estimator of  $\theta$

$$\begin{aligned}Var(\theta) &= E[(\mathbf{b} - \theta)(\mathbf{b} - \theta)^T | X] \\ &= E[(X^T X)^{-1} X^T \epsilon \epsilon^T X (X^T X)^{-1} | X] \\ &= (X^T X)^{-1} X^T E[\epsilon \epsilon^T | X] X (X^T X)^{-1} \\ &= (X^T X)^{-1} \sigma^2\end{aligned}$$

## 3 Implementing linear regression and regularized linear regression(90 points)

### 3.1 Implementing linear regression(45 points)

- A1: Computing the cost function  $J(\theta)$   
See the implementation for the *loss* function in the file *linear\_regressor.py*.
- A2: Implementing gradient descent  
See the implementation for the *train* function in the file *linear\_regressor.py*.




Figure 5: Fitting a linear model




Figure 6: Convergence of gradient descent

- A3: Predicting on unseen data  
For lower status percentage = 5, we predict a median home value of 298034.49  
For lower status percentage = 50, we predict a median home value of -129482.13
- B1: Feature normalization  
See the implementation for the *feature\_normalize* function in the file *utils.py*.
- B2: Loss function and gradient descent  
See the implementation for the *train* and *loss* function in the file *linear\_regressor\_multi.py*.


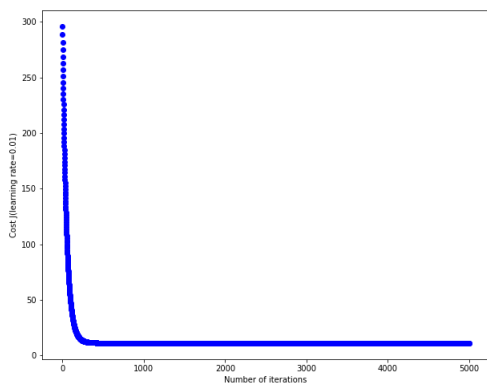



Figure 7: Convergence of gradient descent with multiple variables

- B3: Making predictions on unseen data  
For average home in Boston suburbs, we predict a median home value of 230467.11
- B4: Normal equations  
See the implementation for the *normal\_eqn* function in the file *linear\_regressor\_multi.py*.  
Theta computed by direct solution is:  $[3.64594884e + 01 - 1.08011358e - 014.64204584e - 022.05586264e - 022.68673382e + 00 - 1.77666112e + 013.80986521e + 006.92224640e - 04 - 1.47556685e + 003.06049479e - 01 - 1.23345939e - 02 - 9.52747232e - 019.31168327e - 03 - 5.24758378e - 01]$   
For average home in Boston suburbs, we predict a median home value of 230406.54.

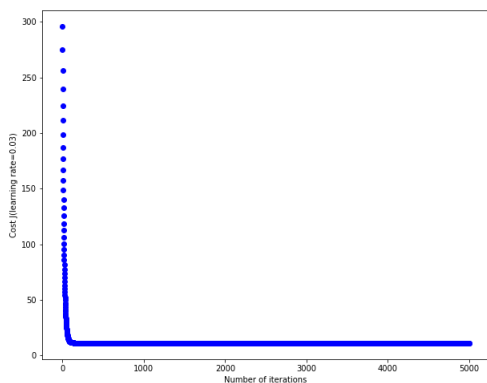
- B5: Exploring convergence of gradient descent




((a)) learning rate=0.01



((b)) learning rate=0.03



((c)) learning rate=0.1

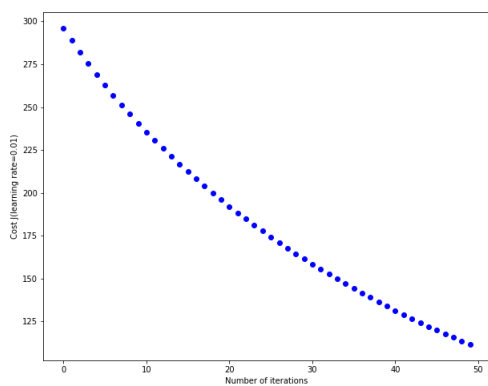


((d)) learning rate=0.3

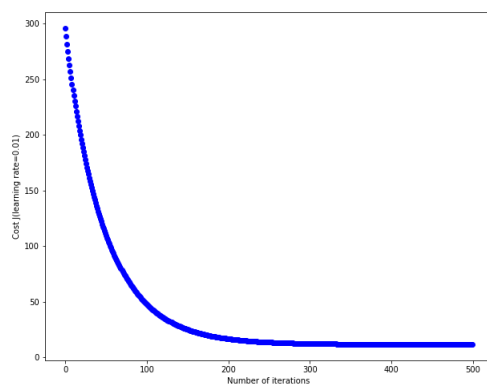
Figure 8: Convergence of gradient descent with multiple variables according to different learning rates

According to the above figures, when learning  $rate = 0.01$  is a good choice.

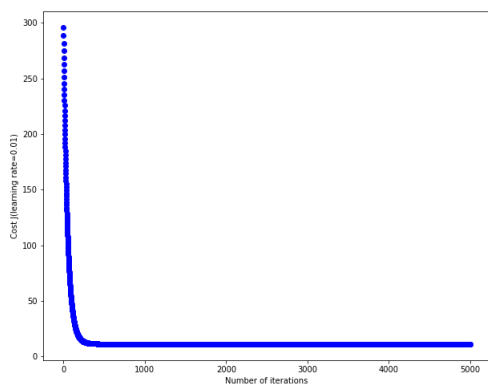
When learning rate = 0.01, choosing different iteration times.



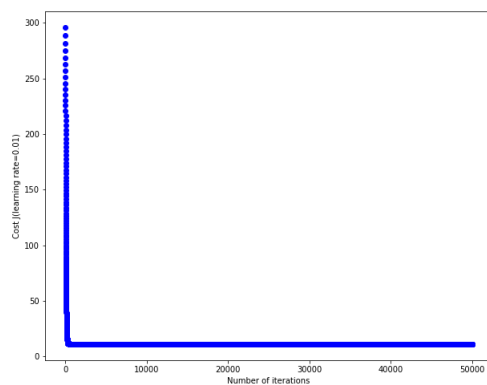
((a)) iteration time = 50



((b)) iteration time = 500



((c)) iteration time = 5000



((d)) iteration time = 50000


Figure 9: Convergence of gradient descent with multiple variables according to different iteration times

According to the above figures, iteration time is 500 is a good choice.

### 3.2 Implementing regularized linear regression(45 points)


- A1: Regularized linear regression cost function  
See the implementation for the *loss* function in the class *Reg\_LinearRegression\_SquaredLoss* in the file *reg\_linear\_regressor\_multi.py*.
- A2: Gradient of the Regularized linear regression cost function  
See the implementation for the *grad\_loss* function in the class *Reg\_LinearRegression\_SquaredLoss* in the file *reg\_linear\_regressor\_multi.py*.
- A3: Learning curves
- A4: Adjusting the regularization parameter

Polynomial Regression fit with lambda = 0.0 and polynomial features of degree 4



((a)) Polynomial Regression fit


Learning curve for linear regression with lambda = 0.0



((b)) Learning curve for linear regression


Figure 10:  $\lambda = 0.0$

Polynomial Regression fit with lambda = 1.0 and polynomial features of degree 4



((a)) Polynomial Regression fit

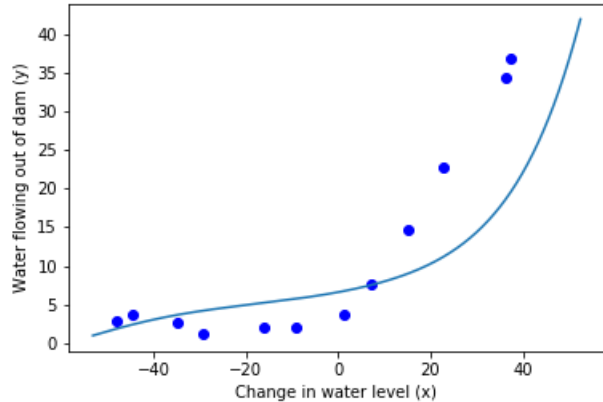
Learning curve for linear regression with lambda = 1.0



((b)) Learning curve for linear regression


Figure 11:  $\lambda = 1.0$

Polynomial Regression fit with lambda = 10.0 and polynomial features of degree 4



((a)) Polynomial Regression fit


Learning curve for linear regression with lambda = 10.0



((b)) Learning curve for linear regression


Figure 12:  $\lambda = 10.0$

Polynomial Regression fit with lambda = 100.0 and polynomial features of degree 4



((a)) Polynomial Regression fit

Learning curve for linear regression with lambda = 100.0



((b)) Learning curve for linear regression

Figure 13:  $\lambda = 100.0$

- A5: Selecting  $\lambda$  using a validation set According to the figure above, the best  $\lambda$  is 3.




Figure 14: Fitting a linear model

- A6: Computing test set error  
Test error is 8.936718220214782.
- A7: Plotting learning curves with randomly selected examples
- A8: Comparing ridge regression and lasso regression models  
Lasso regression converges faster than Ridge regression.

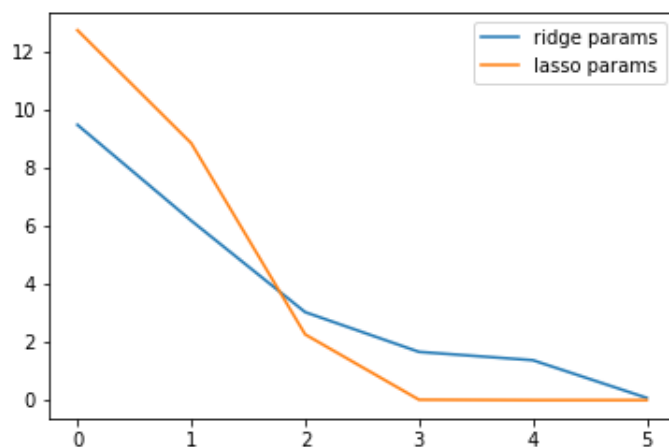


Figure 15: Fitting a linear model

## 4 Extra Credits

- The lowest achievable error on the test set with  $\lambda = 0$   
Test error is 12.930601133114878 when  $\lambda = 0$ .

- Select the best value for  $\lambda$  and report the test set error with the best  $\lambda$ .  
The best  $\lambda = 30.0$ . Test error at  $\lambda = 30.0$  is 13.34301316716897.




Figure 16: Fitting a linear model

- What is the test set error with quadratic features with the best  $\lambda$  chosen with the validation set?  
The best  $\lambda = 30.0$ . Test error at  $\lambda = 30.0$  is 4.742818184669156.




Figure 17: Fitting a linear model



- What is the test set error with cubic features with the best  $\lambda$  chosen with the validation set?  
The best  $\lambda = 300.0$ . Test error at  $\lambda = 30.0$  is 6.558244031190886.




Figure 18: Fitting a linear model